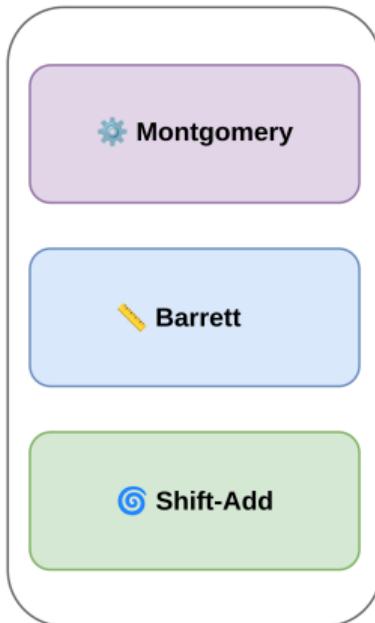


Hardware Evaluation of Modular Multiplication methods

Teodora Alexandrescu

Supervisor: Aikata Aikata

Modular Reduction Algorithms



- Dilithium & Kyber moduli for proof-of-concept
- Area and performance results
- Suggest appropriate implementation strategies

Modular Reduction Algorithms



- Dilithium & Kyber moduli for proof-of-concept
- Area and performance results
- Suggest appropriate implementation strategies

Applications of Modular Multiplication

- Schemes based on LWE (or MLWE, RLWE)

$$\mathbf{A} \cdot \mathbf{s} + \mathbf{e} \approx \mathbf{b} \pmod{q}$$

- Post-quantum schemes: Kyber, Dilithium, Falcon \rightarrow MLWE/ RLWE

Applications of Modular Multiplication

- Schemes based on LWE (or MLWE, RLWE)

$$\mathbf{A} \cdot \mathbf{s} + \mathbf{e} \approx \mathbf{b} \pmod{q}$$

- Post-quantum schemes: Kyber, Dilithium, Falcon \rightarrow MLWE/ RLWE

Why Optimize Modular Reduction?

- Schoolbook reduction: long division by q ($x \bmod q$)
- Expensive operation on digital platforms
- Solution?



Montgomery Reduction

- Montgomery Form
 - $\bar{x} = xR \bmod Q$
- Replace division with shifts by R .

Algorithm 1 Montgomery Reduction (REDC)

Require: $T, Q, R = 2^n$, where $\gcd(R, Q) = 1$.

Output: $TR^{-1} \bmod Q$

```
1:  $m \leftarrow (T \bmod R)Q' \bmod R$ 
2:  $t \leftarrow (T + m \cdot Q)/R$ 
3: if  $t \geq Q$  then
4:    $t \leftarrow T - Q$ 
5: end if
6: return  $t$ 
```

Barrett Reduction

- Approximates q
- Precomputation term μ
- Leverage precomputed term μ to avoid division

Algorithm 3 Barrett Reduction

Require: X, Y, M , $X, Y < M$

Output: $X \cdot Y \bmod M$

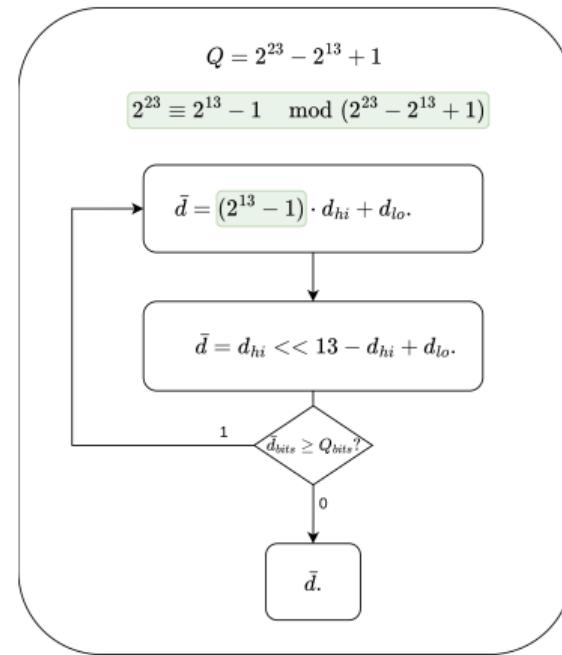
```
1:  $T \leftarrow X \cdot Y$ 
2:  $\mu \leftarrow \lfloor (1/M) \cdot 2^{2k} \rfloor$ 
3:  $q \leftarrow \lfloor (T \cdot \mu) / 2^{2k} \rfloor$ 
4:  $r \leftarrow T - q \cdot M$ 
5: if  $r \geq M$  then
6:    $r \leftarrow r - M$ 
7: end if
8: return  $r$ 
```

- In practical applications, such as lattice-based PQC, often the moduli used are almost powers of two
- General Montgomery and Barrett techniques do not exploit the pseudo-Mersenne form of the moduli

- In practical applications, such as lattice-based PQC, often the moduli used are almost powers of two
- General Montgomery and Barrett techniques do not exploit the pseudo-Mersenne form of the moduli

Shift-Add Reduction

- *Special form modulus*
 - Exploit modular equivalences
 - Compute modulus with *shifts and additions, recursively.*
- Each fold
 - Shrinks the size of the number
 - Maintains congruence modulo q



Can we avoid runtime recursion?

- Precompute recursive approach with pen & paper and *avoid runtime recursion altogether*

$$d[23 : 0] = 2^{12}d[23 : 12] + d[11 : 0]$$

use $2^{12} \equiv 2^9 + 2^8 - 1$.

$$= 2^9d[23 : 12] + 2^8d[23 : 12] - d[23 : 12] + c[11 : 0]$$

...

$$\begin{aligned} &= 2^9(d[23 : 15] + d[23 : 16] + d[14 : 12]) + \\ &\quad 2^8(d[23 : 15] + d[23 : 16] + d[15 : 12]) - \\ &\quad d[23 : 15] - d[23 : 16] - d[23 : 12] + d[11 : 0]. \end{aligned}$$

Precomputed Shift-Add Kyber

$$d[45 : 0] = 2^{23}d[45 : 23] + d[22 : 0]$$

use $2^{23} \equiv 2^{13} - 1$.

$$= (2^{13} - 1)d[45 : 23] + d[22 : 0]$$

...

$$\begin{aligned} &= 2^{13}(d[32 : 23] + d[42 : 33] + d[45 : 43]) - \\ &\quad d[45 : 23] - d[45 : 33] - d[45 : 43]) + d[22 : 0]. \end{aligned}$$

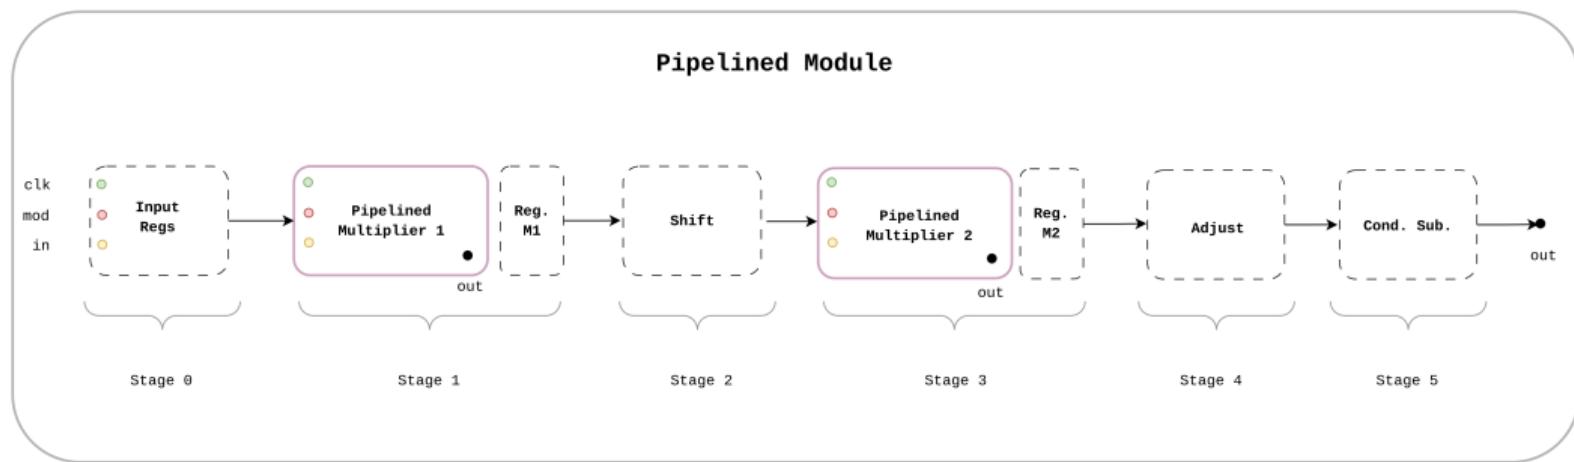
Precomputed Shift-Add Dilithium

Design & Implementation

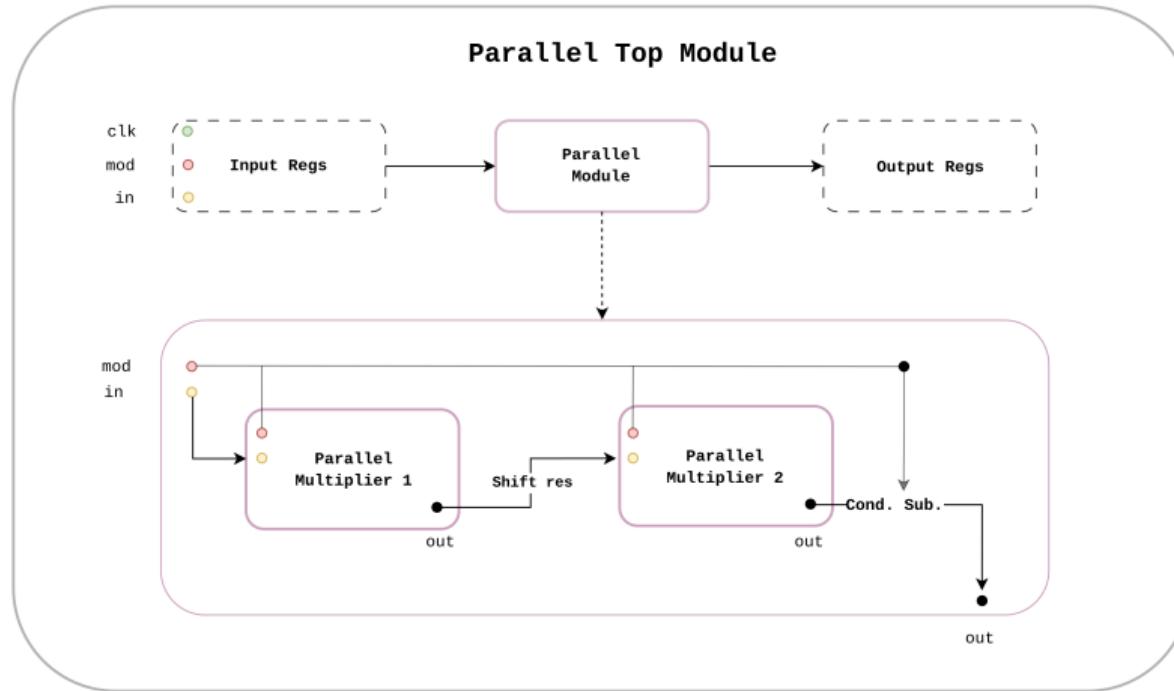
- Approaches and targets
 - *Pipelined*: achieve high throughput
 - *Digit-parallel*: achieve results using few clock cycles
 - *Digit-/Bit-serial*: saves area, but latency is high
- Batch of 64 inputs
- Hardware Description Language: SystemVerilog
- Synthesis Tool: **Xilinx Vivado 2024.2**
- Testing Board: **Artix-7 AC701 Evaluation Platform**

Pipelined

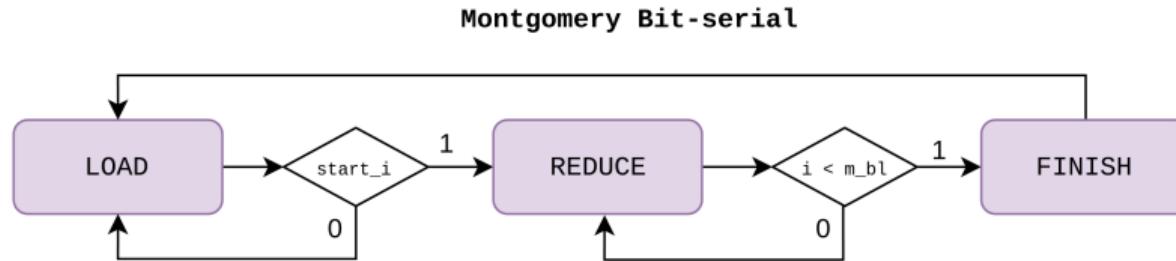
- Approach used by both Barrett & Montgomery



Digit-parallel

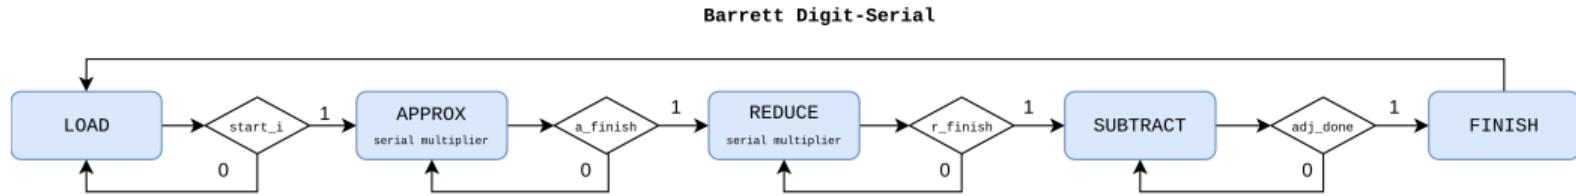


Bit-serial



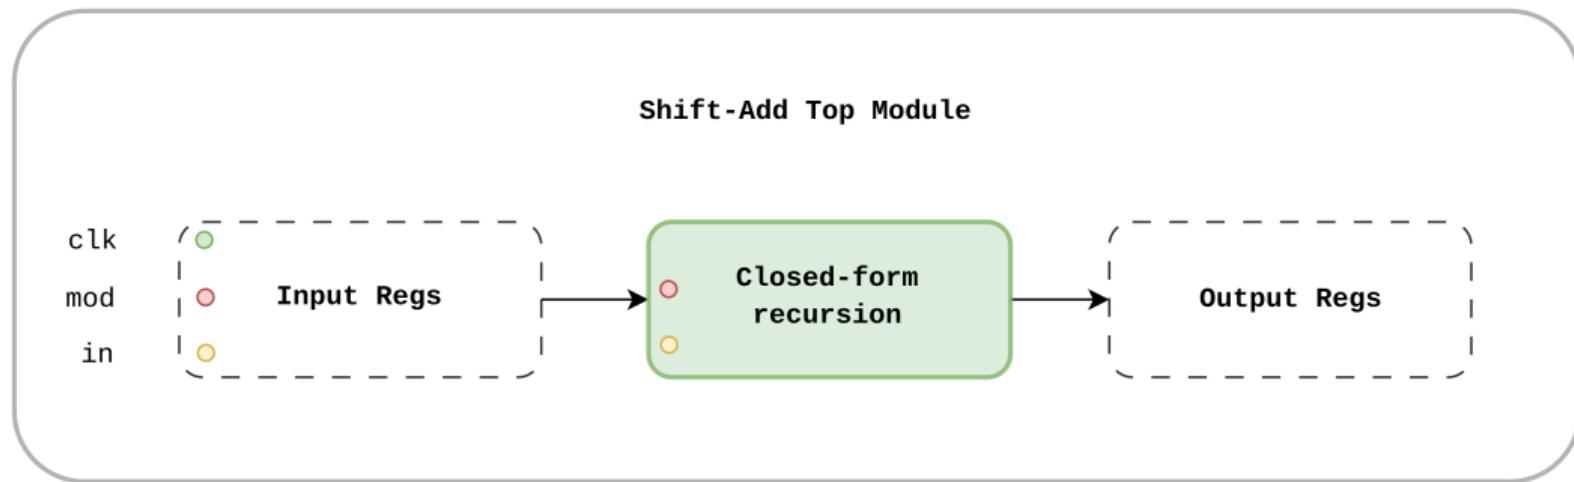
- FSM driven design
- Bit-by-bit processing of an input

Digit-serial



- FSM driven design
- Uses digit-serial modules with *16×16-bit MAC units* for iterative multiplication

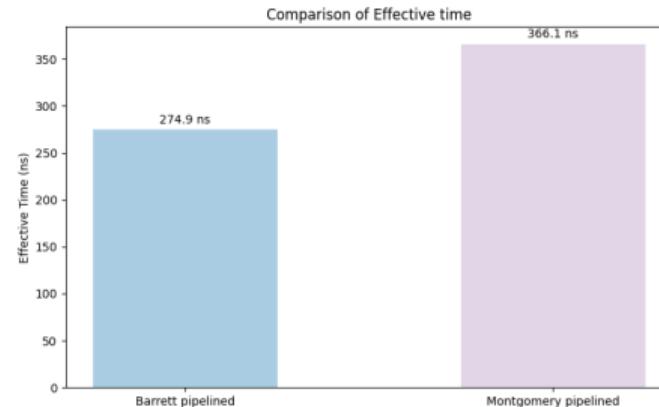
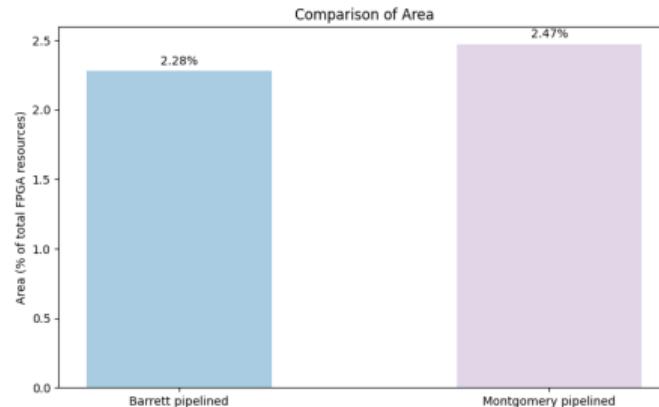
Closed-form Shift-Add



- Implements closed-form logic in one submodule
- Input and output are just registered in the top module

Theory aside - let's put these devices under test!

Pipelined

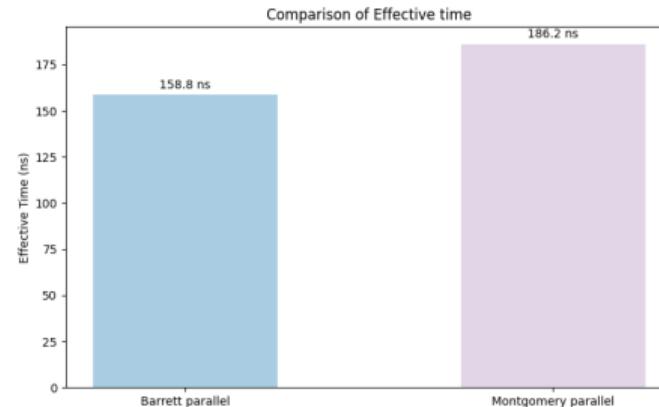
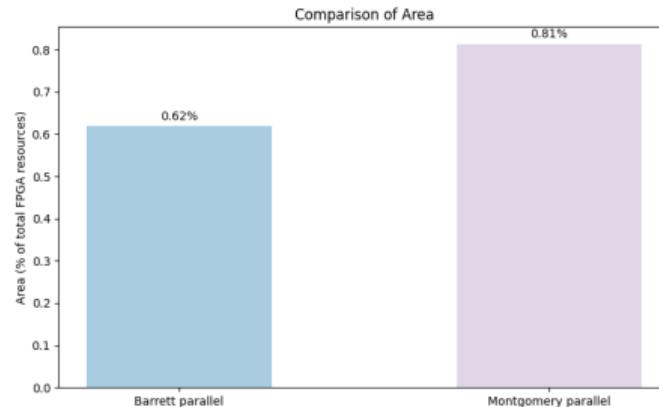


	LUT	FF	DSP
Barrett	1392	2306	26
Montgomery	1606	2441	26

	Frequency (MHz.)	Throughput (outs/s)
Barrett	156.39	152.76
Montgomery	117.46	114.73

Results

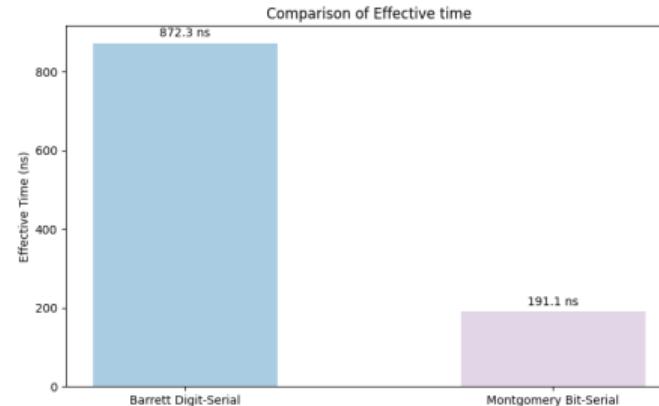
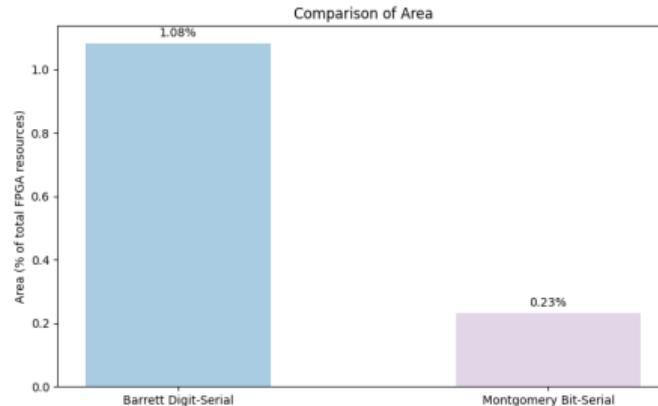
Digit-parallel



	LUT	FF	DSP
Barrett	960	129	26
Montgomery	1606	2441	26

	Frequency (MHz.)	Throughput (outs/s)
Barrett	31.70	6.34
Montgomery	26.85	5.37

Digit-/Bit-Serial

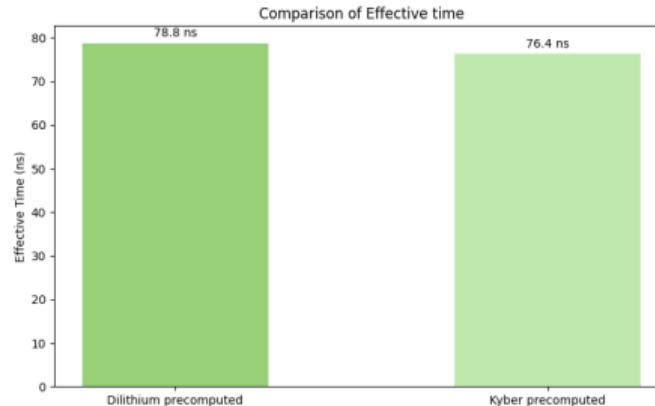
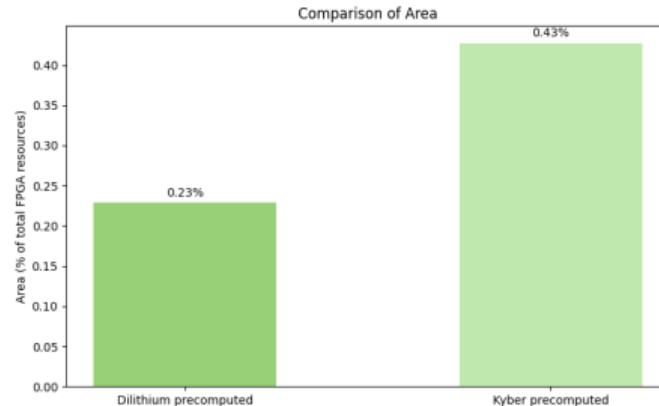


	LUT	FF	DSP
Barrett	948	1008	2
Montgomery	403	75	0

	Frequency (MHz.)	Throughput (outs/s)
Barrett	120.36	1.14
Montgomery	130.80	4.69

Results

Closed-Form Shift-Add



	LUT	FF	DSP
Dilithium	299	111	0
Kyber	631	138	0

	Frequency (MHz.)	Throughput (outs/s)
Dilithium	50.78	10.15
Kyber	53.06	10.61

Conclusion

- Modular multiplication is often the *performance bottleneck* in cryptographic algorithms
- High performance solutions demand further research

Conclusion

- Modular multiplication is often the *performance bottleneck* in cryptographic algorithms
- High performance solutions demand further research

Conclusion

- Modular multiplication is often the *performance bottleneck* in cryptographic algorithms
- High performance solutions demand further research

Hardware Evaluation of Modular Multiplication methods

Teodora Alexandrescu

Supervisor: Aikata Aikata